

Media Engineering and Technology Faculty
German University in Cairo



Bachelor Thesis Title

Bachelor Thesis

Author: Maggie Ezzat Gamil Gaid

Supervisors: Sup 1

Sup 2

Sup 3

Submission Date: XX July, 20XX

Media Engineering and Technology Faculty
German University in Cairo



Bachelor Thesis Title

Bachelor Thesis

Author: Maggie Ezzat Gamil Gaid

Supervisors: Sup 1

Sup 2

Sup 3

Submission Date: XX July, 20XX

This is to certify that:

- (i) the thesis comprises only my original work toward the Bachelor Degree
- (ii) due acknowledgement has been made in the text to all other material used

Maggie Ezzat Gamil Gaid
XX July, 20XX

Acknowledgments

Text

Abstract

Abstact

Contents

Acknowledgments	V
1 Introduction	1
1.1 Section Name	1
1.2 Another Section	1
2 Literature Review	3
2.1 Natural Language Processing	3
2.2 Artificial Neural Networks	3
2.2.1 Feed Forward Neural Networks	4
2.2.2 Recurrent Neural Networks	5
2.2.3 Convolution Neural Networks	8
2.3 Speech Recognition	8
2.3.1 End-to-End Speech Recognition	8
2.3.2 Feature Extraction	8
2.3.3 Building the Model	9
2.3.4 Computing the Loss	9
2.4 Text Analytics	9
2.4.1 Encoder-Decoder Architecture	9
2.4.2 Attention Mechanism	10
2.4.3 The Transformer	11
2.4.4 Bidirectional Encoder from Transformer (BERT)	12
3 Methodology	13
3.1 Datasets	13
3.1.1 Automatic Speech Recognition (ASR) Datasets	13
3.1.2 Text Classifier Datasets	13
4 Results	15
5 Conclusion	17
6 Future Work	19
Appendix	20
A Lists	21
List of Abbreviations	21
List of Figures	22
References	24

Chapter 1

Introduction

1.1 Section Name

1.2 Another Section

Chapter 2

Literature Review

2.1 Natural Language Processing

Natural Languages refer to languages written and spoken by human beings. They are English, Chinese, French, etc. Humans can easily learn and understand such languages. On the other hand, computers have difficulty understanding these languages because of the “ambiguity” problem. Computers understand structured and unambiguous programming languages. After all, on the lowest level everything translates to 0s and 1s.

Natural Language Processing (NLP) is a field which aims at enhancing the human-computer interaction by giving computers the ability to understand natural language. It encompasses two sub-fields: Natural Language Generation (NLG) and Natural Language Understanding (NLU). NLG targets making computers generate human-like sentences. NLU focuses on semantics extraction or building a comprehension of the intent. For quite a long time (since 1950's), NLP researches have been striving to build language models for narrowing down the gap between humans and machines.

In this Literature Review, a brief about two sub domains is included: Speech Recognition and Text Analytics, however, in the beginning a quick review of the widely popular Artificial Neural Networks is elaborated in the next section.

2.2 Artificial Neural Networks

Artificial Neural Network (ANNs) are models of computation modeled after the biological neural networks that constitute the human brain. In early researches of Neural Networks, biological resemblance was emphasized [12] [13] [7], however, nowadays it is obvious that ANNs have little in common compared to biological neural networks. That is due to the biological emphasis being abandoned in favor of achieving satisfying computational results.

The basic building block of ANNs are “neurons”, commonly called nodes. The set of nodes are connected to each other with weighted edges, with the weights of the edges resembling the strength of the “synapses” between the neurons as suggested by the original biological model. The neurons are represented diagrammatically by drawing them as circles, and the weighted edges as arrows connecting them. Each node has an activation function associated with it, which takes as input a weighted sum of the values of input nodes (Figure 2.1).

The most popular activation functions are: the sigmoid function $\sigma(z) = \frac{1}{1+e^{-z}}$, the tanh function $\phi(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$, and the Re-Lu function $l(z) = \max(0, z)$. The activation function at the output nodes is considered to be task-specific. However, the most popular ones are the softmax function $\sigma(\mathbf{z})_i = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}$ for $i = 1, \dots, K$ and $\mathbf{z} = (z_1, \dots, z_K) \in \mathbb{R}^K$, the sigmoid function, or simple linear functions.

There exist many variations of ANNs, the simplest form of them are those whose edges do not form any cycles. These are called Feed Forward Neural Networks.

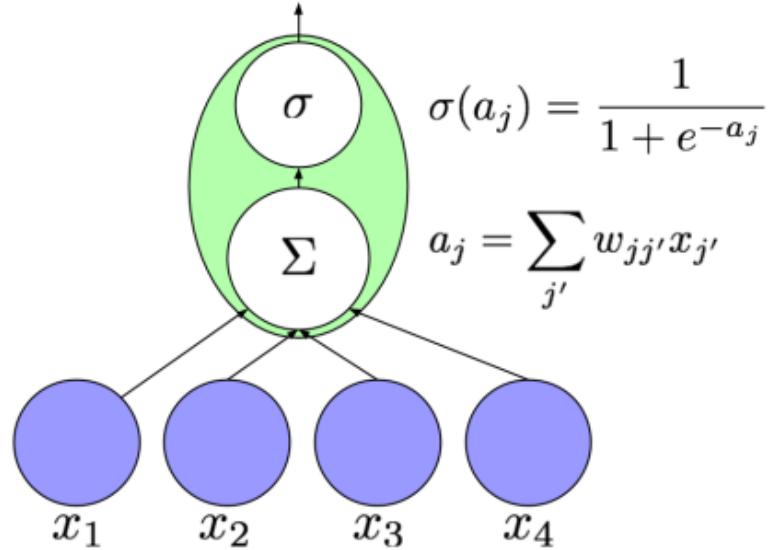


Figure 2.1: A neuron represented as a circle and the weighted edges as arrows. The activation function is a function of the sum of the weighted edges

2.2.1 Feed Forward Neural Networks

The most popular form of Feed Forward Neural Networks (FNNs) is the Multi Layer Perceptrons (MLPs) [14] [17] [4]. With the absence of cycles, the nodes are arranged into layers, as seen in Figure 2.2 and Figure 2.3, where the layers are classified as an input layer, one or many hidden layers, and an output layer.

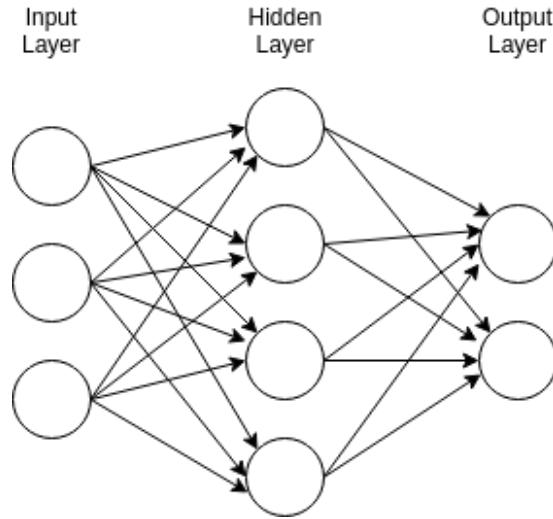


Figure 2.2: A simple Feed Forward Neural Network consisting of an input layer, one hidden layer, and an output layer

The input to FNNs is applied to the input layer. Values of nodes in a given layer, are successively calculated using the values of nodes in the lower layer, until the output is generated at the highest layer: the output layer. This is known as the “Forward pass”. Neural Networks learn by looking at input examples without being explicitly programmed any hard-coded rules about the required task. The learning process is achieved by continuously modifying the weights to minimize an error represented by a loss function $L(\hat{y}, y)$, which measures the distance between the output y (predicted value of y) and the actual value of y (ground-truth).

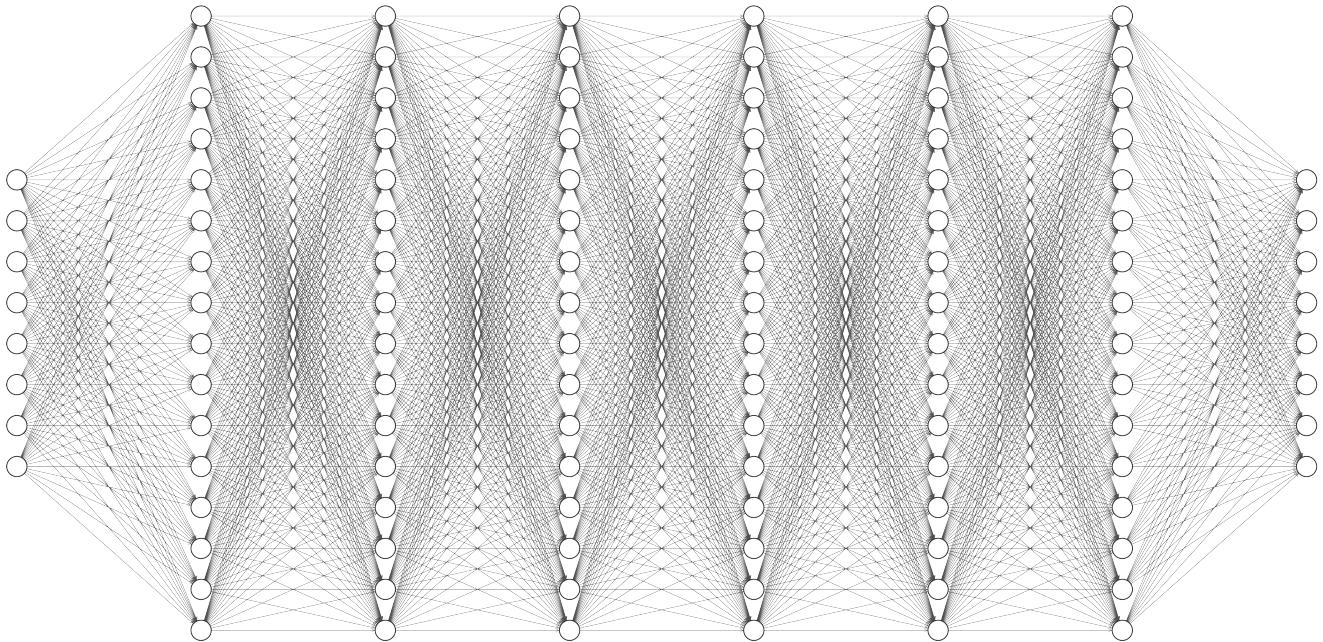


Figure 2.3: A deep Feed Forward Neural Network consisting of many hidden layers

The algorithm for training neural networks is back-propagation [14]. Back-propagation uses the chain rule to calculate the derivative of the loss function $L(\hat{y}, y)$ with respect to each parameter in the network. The parameters (weights) are then adjusted in the direction of less error by an optimization algorithm called gradient descent. This is known as the “Backward Pass”

Sequence Models and the Problem with FNNs

A distinctive feature of the FNNs is the “independence assumption”. That is the presented examples (data points) are assumed to be independent of each other, rendering the FNNs unable to correctly represent input or output sequences with dependencies either in time or space. Examples are words forming sentences, letters forming words, frames of video, snippets of audio clips, DNA sequences, etc. FNNs know no concept of context when analyzing the given examples, they simply are unable to capture dependencies. With the context being a crucial element when analyzing sequences, a simple solution that addresses that matter is the “time-window” solution, i.e. collect the data from either side of the current input into a window. The fact that the range of useful context (either on the left or the right) vary widely from sequence to sequence and in most cases is unknown makes this approach not very efficient. For example, a model trained using a finite-length context window of length n could never be trained to answer the simple question, “what was the data point seen $n + 1$ time steps ago?”

Another problem with FNNs is that they treat inputs and outputs as “fixed-length vectors”. Some representations, such as sentences can not be represented in such way. Some solutions such as “padding” assume a maximum-length for the inputs and/or outputs. Such approach is not a general one. Thus, it was required to extend these powerful and successful models to better suit the vastly crucial sequence models, and that is where the Recurrent Neural Networks came into picture.

2.2.2 Recurrent Neural Networks

ANNs containing cycles are referred to as recursive, or recurrent neural networks. Recurrent Neural Networks (RNNs) are models which have the ability to pass information learnt across past data points, while processing sequential data points one by one. Thus they can model inputs and outputs

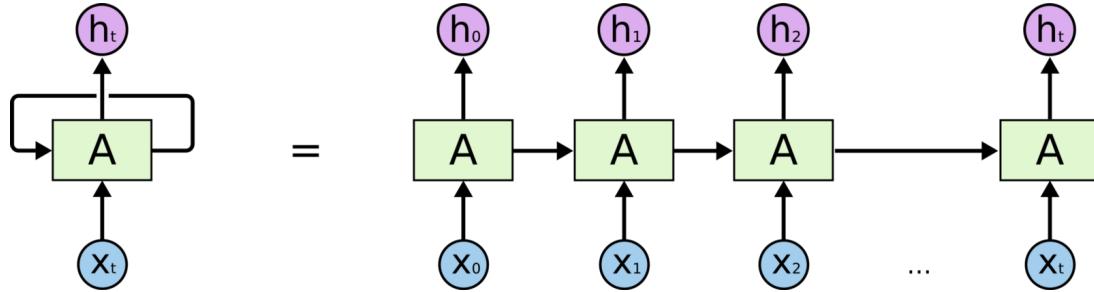


Figure 2.4: A Recurrent Neural Network

which are correlated either in time or space. They are considered to be neural networks possessing memory.

The RNNs have the following architecture: each hidden layer - commonly referred to as “hidden state” - has two sources of inputs, which are the present data point, and information from the hidden state of the past data point (Figure 2.4). This is how contextual information is propagated across the hidden states of the sequential data points. Equation 2.1 explains how each hidden state nodes values are calculated. Each hidden state \mathbf{h}_t is a function of the present data point \mathbf{x} multiplied by some weight matrix W^x and the previous hidden state \mathbf{h}_{t-1} multiplied by some weight matrix W^h and some bias term \mathbf{b}_h . The weight matrices are used to determine how much importance is given to both the present data point and the past hidden state. The output $\hat{\mathbf{y}}_t$ at time step t is given by equation 2.2, where $\hat{\mathbf{y}}_t$ is obtained by applying the softmax function to the hidden state \mathbf{h}_t multiplied by some weight matrix W^y and adding to it some bias term \mathbf{b}_y . Similar to the vanilla ANNs, the weights are continuously adjusted to minimize a cost function. This is done using an algorithm called Backpropagation Through Time (BPTT) [18]

$$\mathbf{h}_t = \phi(W^x \mathbf{x} + W^h \mathbf{h}_{t-1} + \mathbf{b}_h) \quad (2.1)$$

$$\hat{\mathbf{y}}_t = \text{softmax}(W^y \mathbf{h}_t + \mathbf{b}_y) \quad (2.2)$$

Bidirectional RNNs

A slight variation of the RNNs is the Bidirectional Recurrent Neural Networks (BRNNs) [15]. The BRNNs have a slight different architecture which allows it to take into consideration information not only from the present and the past input but also from the future input. Note that by past, present, and future, we not only refer to temporal sequences, but any sequences which have a strong emphasis of the order, but bear no explicit notion of time; this is actually the case with NLP. In BRNNs each hidden layer is duplicated into two layers (Figure 2.5), one layer takes as inputs the present data point and information from the past data point. This is referred to as the “forward direction”. The other layer takes as input the present data point and information from the future data point. This is referred to as the “backward direction”. The BRNNs are fully described by equations 2.3, 2.4 and 2.5, where $\mathbf{h}_t^{<\mathbf{f}>}$ represents the forward direction of the hidden layers and $\mathbf{h}_t^{<\mathbf{b}>}$ represents the backward direction of the hidden layers. The predicted value $\hat{\mathbf{y}}_t$ is now a function of both the forward and the backward direction. Considering information from both sides of the sequence instead of the left side only adds much power to the network as the context is understood much better. Consider the following example: “She said, ‘Teddy bears are on sale.’”, “She said, ‘Teddy Roosevelt was an amazing president’” On these two examples, considering “Teddy” as the current data point, the left sequence is the same, however, the right sequence is crucial in understanding the context. Thus for an application like named-entity recognition, using BRNNs adds much gain. One drawback about BRNNs is that the entire sequence is needed before any predictions can be made, therefore for real-time systems it is a bit slow as it introduces some delay.

$$\mathbf{h}_t^{<\mathbf{f}>} = \phi(W^{xf} \mathbf{x} + W^{hf} \mathbf{h}_{t-1}^{<\mathbf{f}>} + \mathbf{b}_{hf}) \quad (2.3)$$

$$\mathbf{h}_t^{<\mathbf{b}>} = \phi(W^{xb} \mathbf{x} + W^{hb} \mathbf{h}_{t+1}^{<\mathbf{b}>} + \mathbf{b}_{hb}) \quad (2.4)$$

$$\hat{\mathbf{y}}_t = \text{softmax}(W^{yf} \mathbf{h}_t^{\mathbf{f}} + W^{yb} \mathbf{h}_t^{\mathbf{b}} + \mathbf{b}_y) \quad (2.5)$$

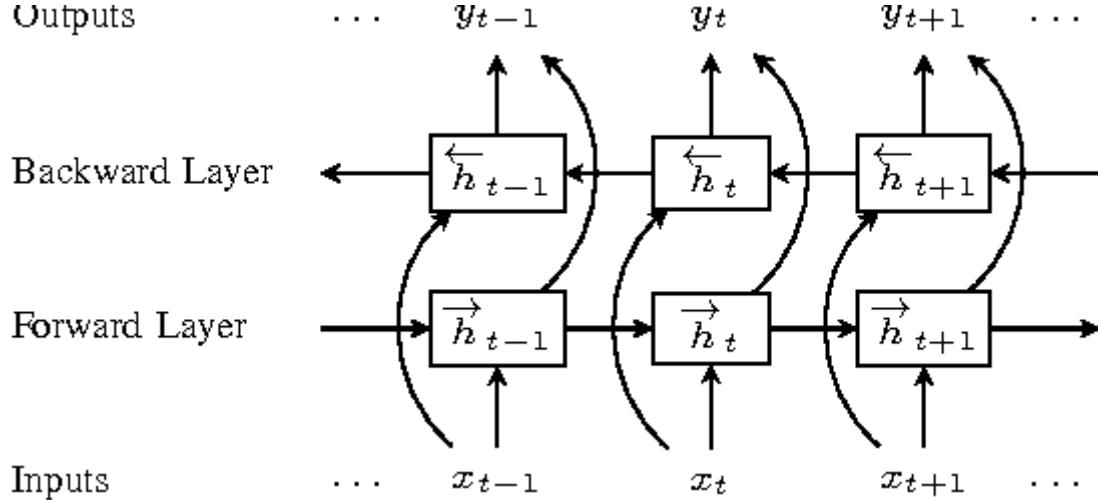


Figure 2.5: A Bidirectional Recurrent Neural Network

Problems with RNNs

TODO: vanishing and exploding gradients [9] [10] [3]

Long Short-Term Memory

Long Short-Term Memory (LSTM) [11] is a variation of vanilla RNNs which was designed as a solution to the vanishing gradients problem. The main idea was to replace the ordinary node with a “memory cell”. The cell uses “gates” in order to make decisions about which information to keep and which to discard. It has three gates: output, input, and forget gate, which are analogous to read, write, and reset operations for the memory cell. The gates uses sigmoid as the activation function, where the values of the gates ranges from 0 to 1.

TODO: continue LSTM

$$\tilde{\mathbf{c}}_t = \tanh(W^{ch} \mathbf{h}_{t-1} + W^{cx} \mathbf{x}_t + \mathbf{b}_c) \quad (2.6)$$

$$\mathbf{f}_t = \sigma(W^{fh} \mathbf{h}_{t-1} + W^{fx} \mathbf{x}_t + \mathbf{b}_f) \quad (2.7)$$

$$\mathbf{i}_t = \sigma(W^{ih} \mathbf{h}_{t-1} + W^{ix} \mathbf{x}_t + \mathbf{b}_i) \quad (2.8)$$

$$\mathbf{o}_t = \sigma(W^{oh} \mathbf{h}_{t-1} + W^{ox} \mathbf{x}_t + \mathbf{b}_o) \quad (2.9)$$

$$\mathbf{c}_t = \mathbf{i}_t \tilde{\mathbf{c}}_t + \mathbf{f}_t \mathbf{c}_{t-1} \quad (2.10)$$

$$\mathbf{h}_t = \tanh(\mathbf{c}_t) \mathbf{o}_t \quad (2.11)$$

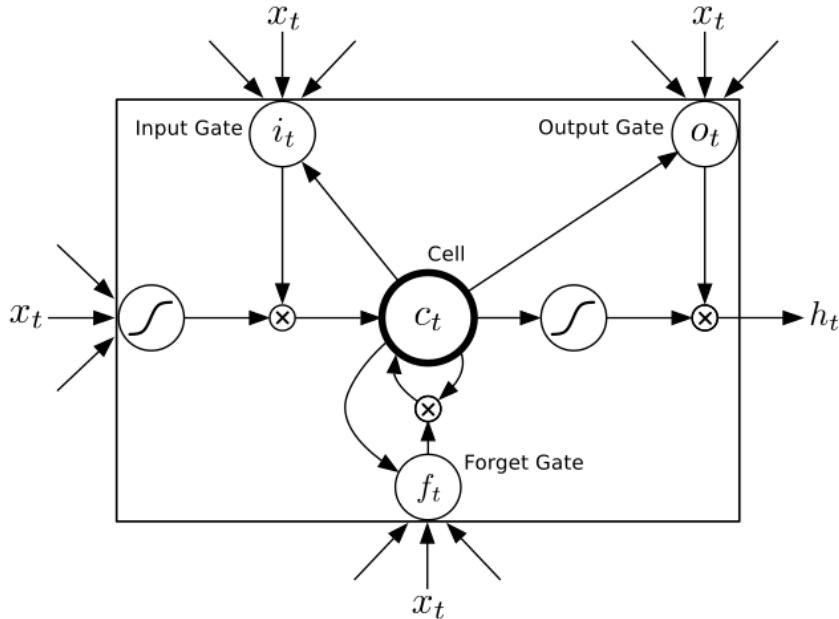


Figure 2.6: An LSTM Cell

Gated Recurrent Units

TODO: Gated Recurrent Units (GRUs)

2.2.3 Convolution Neural Networks

2.3 Speech Recognition

2.3.1 End-to-End Speech Recognition

2.3.2 Feature Extraction

The premier step in any ASR system is to extract features. That is to turn the raw speech waveforms into a sequence of numerical vectors. The issue, however, lies in the fact that the audio signals are constantly changing. For us to be able to deal with them, we make the assumption that for sufficiently small period of time the signals are stationary. Therefore we sample the signal into 25ms frames, with a step of about 10ms which so that the frames are overlapping. Then we apply some operations on each frame to extract features. The most widely used feature extraction method is Mel-Frequency Cepstral Coefficients (MFCC), which we explain in brief.

For each frame, we calculate the Fast Fourier Transform (FFT), in order to move from the time domain to the spectral domain. Human ears cannot distinguish between two closely spaced frequencies, specially for higher frequencies. To mimic this effect, we need to know how much energy exists in different frequency regions. This is done by applying special filters called Mel filterbanks. The first filter is the narrowest, and indicates how much energy exists around 0 Hertz. Then the filters become wider as the frequencies get higher because we become less concerned about variations. Afterward, we take the log of the filterbank energies. This is also done to mimic the human ear as we do not hear loudness on a linear scale. Then we compute the Discrete Cosine Transform (DCT) of the log filterbank energies. This is due to the fact that the filterbanks are overlapping, so we compute DCT to reduce the correlation between filter bank amplitudes. The resulting DCT coefficients are referred

to as MFCC coefficients. Only 12 coefficients are kept and the rest are dropped, this is proven to improve the ASR performance. These 12 coefficients, together with the normalized energy, they form the feature vector.

2.3.3 Building the Model

2.3.4 Computing the Loss

2.4 Text Analytics

Text Analytics is the process of extracting information and semantics from written text. Some famous text-analytics tasks are text classification, sentiment analysis, document summarization, named-entity recognition and entity-relation extraction. Due to the sequential nature of text, Recurrent Neural Network (RNN)s were widely used for many text analytics tasks for quite a long time. This was the case because their sequential nature made them a good candidate for the task. A special architecture, named “Encoder-Decoder Architecture”, which made use of RNNs when first introduced, has been marked broadly useful in modeling sequence-to-sequence models, where the input is a sequence, and the output is as well a sequence. An example is machine translation, where a model is trained to find an output sentence y which maximizes the conditional probability of y given an input sentence x . We discussed RNNs and their variations in section 2.2.2. In the following section, we take a look into the encoder-decoder architecture.

2.4.1 Encoder-Decoder Architecture

The popular encoder-decoder architecture which was first proposed by Cho *et al.* (2014a) [6] and Sutskever *et al.* (2014) [16]. This system consists of two RNNs which work together as an encoder-decoder pair. The encoder encodes a variable-length sequence (e.g. a sentence) into a fixed-length vector which we call summary vector \mathbf{c} . The decoder then uses this fixed-length vector to generate a variable-length output sequence. The vector \mathbf{c} has information from each data point in the input sequence.

The encoder is a RNN which has a hidden state updated at each time step according to equation 2.12, where f is a non-linear activation function; Sutskever *et al.* (2014) [16] uses LSTMs for this purpose while Cho *et al.* (2014a) [6] used a variation of LSTMs instead.

$$\mathbf{h}_t = f(x_t, \mathbf{h}_{t-1}) \quad (2.12)$$

The decoder is a RNN that is trained to generate the output sequence \mathbf{y} one by one. It uses the hidden state \mathbf{h}_t to calculate y_t . The hidden state \mathbf{h}_t is calculated according to equation 2.13, where it uses the previously generated symbol y_{t-1} , the previous hidden state \mathbf{h}_{t-1} and the summary vector \mathbf{c} in predicting the target symbol. Similarly, the conditional probability of the target symbol is given by 2.14

$$\mathbf{h}_t = f(\mathbf{h}_{t-1}, y_{t-1}, \mathbf{c}) \quad (2.13)$$

$$P(y_t | y_1, y_2, \dots, y_{t-1}, \mathbf{c}) = g(\mathbf{h}_t, y_{t-1}, \mathbf{c}) \quad (2.14)$$

The two components are then trained to maximize the conditional probability of the output sequence $\mathbf{y} = (y_1, y_2, \dots, y_{T_y})$ given the input sequence $\mathbf{x} = (x_1, x_2, \dots, x_{T_x})$

The problem with this architecture is that the performance deteriorates as the length of the input sequences increases [5]. This is due to the difficulty of cramming all the necessary information into a fixed-length vector. In order to address this issue, the attention mechanism was introduced by Dzmitry *et al.* [2]

2.4.2 Attention Mechanism

Attention is the ability to focus on important details and discard unimportant or irrelevant information. Dzmitry *et al.* (2015) [2] proposes modifying the encoder-decoder architecture through arming the decoder with attention mechanism which allows it to attend to only parts in the input sentence which are most relevant to the target word in the output sequence. The characteristic feature of this approach is that it doesn't encode all the input sequence into a fixed-length vector as the basic encoder-decoder approach explained in section 2.4.1. Instead, it encodes the input sequence into a number of vectors, and chooses which of these vectors are relevant to the target word in order to make a prediction. This approach performs better on longer input sequences as it is no longer needed to suppress all the information given in the sentence in one fixed-length vector. We demonstrate the model proposed by Dzmitry *et al.* (2015) [2] starting by the encoder, then the decoder.

I. Encoder

No major modifications were performed on the encoder, however, a Bidirectional Recurrent Neural Network (BRNN) was used instead of a uni-directional one. This is done in order to gather left and right context as discussed in section 2.2.2. By concatenating the forward hidden state \vec{h}_j and the backward hidden state \overleftarrow{h}_j for each word x_j in the sequence, an annotation $h_j = [\vec{h}_j^T, \overleftarrow{h}_j^T]$ for that word is obtained. The decoder would use these annotations later in the attention layer as we will explain.

II. Decoder

The decoder is a RNN that also uses the hidden state \mathbf{h}_i to calculate y_i . The hidden state \mathbf{h}_i of the decoder is calculated according to equation 2.15, where it uses the previously generated symbol y_{i-1} , the previous hidden state \mathbf{h}_{i-1} and a context vector \mathbf{c}_i in predicting the target symbol. The conditional probability of the target symbol is given by 2.16. The major difference here from the basic encoder-decoder approach is that there is a distinct context vector \mathbf{c}_i for each target word y_i . The context vector \mathbf{c}_i is calculated as a weighted sum of the sequence of annotations produced by the encoder as seen in equation 2.17. The weights of the annotations are learned by the attention model as in equation 2.18 and 2.19

$$\mathbf{h}_i = f(\mathbf{h}_{i-1}, y_{i-1}, \mathbf{c}_i) \quad (2.15)$$

$$P(y_i | y_1, y_2, \dots, y_{i-1}, \mathbf{x}) = g(\mathbf{h}_i, y_{i-1}, \mathbf{c}_i) \quad (2.16)$$

$$\mathbf{c}_i = \sum_{j=1}^{T_x} \alpha_{ij} h_j \quad (2.17)$$

$$\alpha_{ij} = \frac{\exp(e_{ij})}{\sum_{k=1}^{T_x} \exp(e_{ik})} \quad (2.18)$$

where

$$e_{ij} = a(h_{i-1}, s_i) \quad (2.19)$$

III. Additive Attention

The attention model helps solving long range dependencies efficiently because it can focus on what matter and disregard what doesn't matter.

In 2017, Google came up with a novel architecture using only the attention mechanism and eliminated the use of RNNs, this was called “The Transformer” and we explain the idea behind it in the next section.

2.4.3 The Transformer

I. Problem with Recurrence

As seen in equation 2.11, RNNs calculate the hidden state as a function of the current data point and the previous hidden state. This sequential nature of calculation largely suits the sequential nature of natural languages, however, this introduces a major problem as it hinders parallelization among the training inputs. This issue is manifested when the input sequences are of longer lengths. The Transformer abandons recurrence in order to achieve more parallelization and efficiency in performance.

II. Encoder

The Transformer follows the prominent encoder-decoder architecture, however, introducing major modifications to both the encoder and the decoder. The encoder consists of N layers, each layer made up of two sub-layers. The first sub-layer is called “Self-Attention” or “Intra-Attention” sub-layer. This is one of the most influential changes proposed in the Transformer. What self-attention does is that it calculates the relevance of each word in the input sequence to every other word in the sequence. To demonstrate the gain we achieve with self-attention, consider the following examples which highlight the problem of linking the pronouns to their antecedents:

1. “The animal did not cross the street because *it* was too tired.”

When translated to French this becomes: “L’animal n’a pas traversé la rue car *il* était très fatigué”.

2. “The animal did not cross the street because *it* was too wide.”

In French, this is “L’animal n’a pas traversé la rue car *elle* était très large.”

A model implementing no self-attention mechanism has difficulty determining the right pronoun in French; whether “il” for males or “elle” for females, due to the fact that it did not learn any link between “it” and its antecedent when encoding the input sentence.

The second sub-layer is a simple fully-connected feed-forward neural network. A residual connection [8] is applied to each sub-layer, then layer normalization [1] is performed.

III. Decoder

The decoder is made up from N layers as well, with each layer composed of three sub-layers: a self-attention layer similar to the self-attention layer implemented in the encoder, however, modified so that each word can attend only to words earlier in the sequence and not to consequent words, a feed-forward network, and an additional layer which implements the encoder-decoder attention as explained in section 2.4.2. As in the encoder, residual connections are applied on each sub-layer followed by layer normalization.

IV. Scaled Dot-Product Attention

V. Positional Encodings

2.4.4 Bidirectional Encoder from Transformer (BERT)

Word Embeddings

Transfer Learning

Pre-Training

Feature Extraction

Fine Tuning

Bidirectional vs uni-directional

PreTraining Tasks

Chapter 3

Methodology

3.1 Datasets

3.1.1 ASR Datasets

End-to-End ASR systems require large amounts of transcribed audio data. For this purpose we make use of three open-source German datasets and clean them. We list the three datasets here with the cleaning operations performed for each.

I. Common Voice

Common Voice is the largest open source, multi-language dataset of voices available for use. It is managed by Mozilla and was collected by volunteers on-line who were either recording samples or validating other samples. Mozilla began work on this project in 2017 and contribution to the dataset continues up till now. For our ASR, the German subset of the dataset was selected. It incorporates 340 total hours, with 325 validated hours and 15 invalidated hours which we excluded. The dataset has 5007 speakers but as we discarded the invalidated utterances we end up with only 4800 speakers. All the utterances were in “mp3” format and sampled using a sampling rate of $44KHz$ so we converted them to “wav” format and we performed down-sampling to obtain sample rate of $16KHz$. We checked for any corrupted files but there were none.

II. M-AILABS

III. German Speech Data - Tuda-De

3.1.2 Text Classifier Datasets

I. German Wikipedia Dump

II. 10K German Articles

Chapter 4

Results

Results

Chapter 5

Conclusion

Conclusion

Chapter 6

Future Work

Text

Appendix

Appendix A

Lists

NLP	Natural Language Processing
NLG	Natural Language Generation
NLU	Natural Language Understanding
ANNs	Artificial Neural Network
FNNs	Feed Forward Neural Networks
MLPs	Multi Layer Perceptrons
RNNs	Recurrent Neural Networks
RNN	Recurrent Neural Network
BPTT	Backpropagation Through Time
BRNNs	Bidirectional Recurrent Neural Networks
BRNN	Bidirectional Recurrent Neural Network
LSTM	Long Short-Term Memory
GRUs	Gated Recurrent Units
ASR	Automatic Speech Recognition
MFCC	Mel-Frequency Cepstral Coefficients
FFT	Fast Fourier Transform
DCT	Discrete Cosine Transform

List of Figures

2.1	A neuron represented as a circle and the weighted edges as arrows. The activation function is a function of the sum of the weighted edges	4
2.2	A simple Feed Forward Neural Network consisting of an input layer, one hidden layer, and an output layer	4
2.3	A deep Feed Forward Neural Network consisting of many hidden layers	5
2.4	A Recurrent Neural Network	6
2.5	A Bidirectional Recurrent Neural Network	7
2.6	An LSTM Cell	8

Bibliography

- [1] J Ba, J Kiros, and G Hinton. Layer normalization. arxiv. 2016.
- [2] Dzmitry Bahdanau, Kyunghyun Cho, and Yoshua Bengio. Neural machine translation by jointly learning to align and translate. *arXiv preprint arXiv:1409.0473*, 2014.
- [3] Yoshua Bengio, Patrice Simard, Paolo Frasconi, et al. Learning long-term dependencies with gradient descent is difficult. *IEEE transactions on neural networks*, 5(2):157–166, 1994.
- [4] Christopher M Bishop et al. *Neural networks for pattern recognition*. Oxford university press, 1995.
- [5] Kyunghyun Cho, Bart Van Merriënboer, Dzmitry Bahdanau, and Yoshua Bengio. On the properties of neural machine translation: Encoder-decoder approaches. *arXiv preprint arXiv:1409.1259*, 2014.
- [6] Kyunghyun Cho, Bart Van Merriënboer, Caglar Gulcehre, Dzmitry Bahdanau, Fethi Bougares, Holger Schwenk, and Yoshua Bengio. Learning phrase representations using rnn encoder-decoder for statistical machine translation. *arXiv preprint arXiv:1406.1078*, 2014.
- [7] Jeffrey L Elman. Finding structure in time. *Cognitive science*, 14(2):179–211, 1990.
- [8] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 770–778, 2016.
- [9] Sepp Hochreiter. Untersuchungen zu dynamischen neuronalen netzen. *Diploma, Technische Universität München*, 91(1), 1991.
- [10] Sepp Hochreiter, Yoshua Bengio, Paolo Frasconi, Jürgen Schmidhuber, et al. Gradient flow in recurrent nets: the difficulty of learning long-term dependencies, 2001.
- [11] Sepp Hochreiter and Jürgen Schmidhuber. Long short-term memory. *Neural computation*, 9(8):1735–1780, 1997.
- [12] John J Hopfield. Neural networks and physical systems with emergent collective computational abilities. *Proceedings of the national academy of sciences*, 79(8):2554–2558, 1982.
- [13] MI Jordan. Serial order: a parallel distributed processing approach. technical report, june 1985–march 1986. Technical report, California Univ., San Diego, La Jolla (USA). Inst. for Cognitive Science, 1986.
- [14] David E Rumelhart, Geoffrey E Hinton, and Ronald J Williams. Learning internal representations by error propagation. Technical report, California Univ San Diego La Jolla Inst for Cognitive Science, 1985.
- [15] Mike Schuster and Kuldip K Paliwal. Bidirectional recurrent neural networks. *IEEE Transactions on Signal Processing*, 45(11):2673–2681, 1997.

- [16] Ilya Sutskever, Oriol Vinyals, and Quoc V Le. Sequence to sequence learning with neural networks. In *Advances in neural information processing systems*, pages 3104–3112, 2014.
- [17] Paul J Werbos. Generalization of backpropagation with application to a recurrent gas market model. *Neural networks*, 1(4):339–356, 1988.
- [18] Paul J Werbos et al. Backpropagation through time: what it does and how to do it. *Proceedings of the IEEE*, 78(10):1550–1560, 1990.